

Lab 10: mmap

1950787 杨鑫

1. 实验目的：

- 了解 mmap 系统调用和 munmap 系统调用的功能和实现方法
- 修改一些数据结构，添加部分代码以实现一个简化版的 mmap 系统调用和 munmap 系统调用
- 了解文件映射等相关操作
- 进一步熟悉懒加载和文件操作

2. 实验步骤：

首先切换分支至 mmap branch，以获取本次实验的内容

同时清理文件，以获得纯净的初始文件系统

```
git checkout mmap
make clean
```

a. mmap

实验目的：

mmap 和 munmap 系统调用允许 UNIX 程序对其地址空间进行详细控制。它们可用于在进程之间共享内存，将文件映射到进程地址空间，以及作为用户级页面错误方案的一部分，例如讲座中讨论的垃圾收集算法。在本实验中，您将向 xv6 添加 mmap 和 munmap，重点关注内存映射文件。

mmap 可以通过多种方式调用，但本实验只需要其与内存映射文件相关的部分功能。你可以假设 addr 总是为零，这意味着内核应该决定映射文件的虚拟地址。mmap 返回该地址，如果失败则返回 0xffffffffffffff。length 是要映射的字节数；它可能与文件的长度不同。prot 指示内存是否应该被映射为可读、可写和/或可执行；您可以假设 prot 是 PROT_READ 或 PROT_WRITE 或两者兼而有之。flags 要么是 MAP_SHARED，这意味着对映射内存的修改应该写回文件，要么是 MAP_PRIVATE，这意味着它们不应该。您不必在标志中实现任何其他位。fd 是要映射的文件的打开文件描述符。您可以假设偏移量为零（它是文件中映射的起点）。

如果映射同一个 MAP_SHARED 文件的进程不共享物理页面，那也没关系。

munmap(addr, length) 应该删除指定地址范围内的 mmap 映射。如果进程修改了内存并映射了 MAP_SHARED，则应首先将修改写入文件。一个 munmap 调用可能只覆盖 mmap-ed 区域的一部分，但您可以假设它会在开始、结束或整个区域取消映射（但不会在区域中间打孔）。

您应该实现足够的 mmap 和 munmap 功能以使 mmptest 测试程序正常工作。

实验分析：

本次实验要求实现将文件映射到用户的虚拟内存区域的 mmap 和 munmap 系统调用。首先类似于第二节系统调用实验的内容，添加 mmap 和 munmap 的声明以及系统调用编号等，以便用户能够使用该系统调用；然后添加虚拟内存区域的数据结构，用于记录其地址，长度等信息，并且添加相关的初始化，分配虚拟内存空间，释放虚拟内存空间的函数；然后完成 mmap 和 munmap 系统调用的函数功能实现：mmap 取出一块空闲的虚拟内存空间，然后修改其长度，权限等信息，并且将文件指针赋给它；munmap 则找到要释放的虚拟内存，如果是被共享的要先将内容写回物理空间，然后再看是否将所有的页面均删除了，如果是的话还要将文件引用 -1，并将该虚拟空间的指针置空，若未删完，则修改其起始地址和长度。

此外，还需要修改 usertrap，因为虚拟内存区域也是采用懒分配的方式所以当产生页错误时，要及时将相应的虚拟内存放入用户的页表中。最后修改 exec 和 fork 系统调用，退出进程的时候将相应虚拟内存释放掉，产生子进程的时候将虚拟内存复制给子进程。

实验核心代码：

vma.h

```
struct vma {
    char* addr;
    uint64 len;
    struct file *file;
    char prot;
    char flag;
};
```

vma.c

```
#include "types.h"
#include "riscv.h"
#include "defs.h"
#include "param.h"
#include "fs.h"
#include "spinlock.h"
#include "sleeplock.h"
#include "file.h"
#include "vma.h"

struct {
    struct spinlock lock;
    struct vma areas[NOFILE];
} vma_table;

void
vma_init(void)
{
    initlock(&vma_table.lock, "vma_table");
}

struct vma*
vma_alloc(void)
{
    struct vma *vmap;

    acquire(&vma_table.lock);
    for(vmap = vma_table.areas; vmap < vma_table.areas + NOFILE; vmap++){
```

```

        if(vmap->file == 0){
            release(&vma_table.lock);
            return vmap;
        }
    }
    release(&vma_table.lock);
    return 0;
}

void
vma_free(struct vma *vmap)
{
    vmap->file = 0;
}

```

sysfile.c

```

uint64 sys_mmap(void){
    int length, prot, flags, fd;
    if (argint(1, &length) < 0 || argint(2, &prot) < 0 || argint(3, &flags) <
0 || argint(4, &fd) < 0){
        return 0xffffffffffffffff;
    }
    struct proc *pr = myproc();
    if(!pr->ofile[fd]->readable){
        if (prot & PROT_READ)
            return 0xffffffffffffffff;
    }
    if(!pr->ofile[fd]->writable){
        if (prot & PROT_WRITE && flags==MAP_SHARED)
            return 0xffffffffffffffff;
    }

    struct vma *vmap;
    if ((vmap = vma_alloc()) == 0){
        return 0xffffffffffffffff;
    }

    acquire(&pr->lock);
    int i;
    for (i = 0; i < NOFILE; i++){
        if(pr->vmmaps[i] == 0){
            pr->vmmaps[i] = vmap;
            release(&pr->lock);
            break;
        }
    }
    if (i == NOFILE){
        return 0xffffffffffffffff;
    }

    uint64 sz = pr->sz;
    pr->sz += length;
    vmap->addr = (char*)sz;
    vmap->len = length;
    vmap->prot = (prot & PROT_READ) | (prot & PROT_WRITE);
    vmap->flag = flags;
}

```

```

vmap->file = pr->ofile[fd];
filedup(pr->ofile[fd]);
return sz;
}

uint64 sys_munmap(void){
    struct proc *pr = myproc();
    int startAddr, length;

    if (argint(0, &startAddr) < 0 || argint(1, &length) < 0){
        return -1;
    }

    for(int i = 0; i < NOFILE; i++){
        if (pr->vmmaps[i] == 0) {
            continue;
        }
        if ((uint64)pr->vmmaps[i]->addr == startAddr){
            if (length >= pr->vmmaps[i]->len) {
                length = pr->vmmaps[i]->len;
            }
            if (pr->vmmaps[i]->prot & PROT_WRITE && pr->vmmaps[i]->flag ==
MAP_SHARED) {
                begin_op();
                ilock(pr->vmmaps[i]->file->ip);
                writei(pr->vmmaps[i]->file->ip, 1, (uint64)startAddr, 0, length);
                iunlock(pr->vmmaps[i]->file->ip);
                end_op();
            }
            uvmunmap(pr->pagetable, (uint64)startAddr, length/PGSIZE, 1);
            if (length == pr->vmmaps[i]->len){
                fileclose(pr->vmmaps[i]->file);
                vma_free(pr->vmmaps[i]);
                pr->vmmaps[i] = 0;
                return 0;
            } else {
                pr->vmmaps[i]->addr += length;
                pr->vmmaps[i]->len -= length;
                return 0;
            }
        }
    }
    return -1;
}

```

trap.c

```

//
// handle an interrupt, exception, or system call from user space.
// called from trampoline.S
//
void
usertrap(void)
{
    int which_dev = 0;

    if((r_sstatus() & SSTATUS_SPP) != 0)

```

```

panic("usertrap: not from user mode");

// send interrupts and exceptions to kerneltrap(),
// since we're now in the kernel.
w_stvec((uint64)kernelvec);

struct proc *p = myproc();

// save user program counter.
p->trapframe->epc = r_sepc();

if(r_scause() == 8){
    // system call

    if(p->killed)
        exit(-1);

    // sepc points to the ecalls instruction,
    // but we want to return to the next instruction.
    p->trapframe->epc += 4;

    // an interrupt will change sstatus & c registers,
    // so don't enable until done with those registers.
    intr_on();

    syscall();
} else if((which_dev = devintr()) != 0){
    // ok
} else if(r_scause() == 13 || r_scause() == 15) {
    uint64 stval = r_stval();
    if (stval >= p->sz) {
        p->killed = 1;
    } else {
        uint64 protectTop = PGROUNDDOWN(p->trapframe->sp);
        uint64 stvalTop = PGROUNDUP(stval);
        if (protectTop != stvalTop) {
            struct vma *vmap;
            int i;
            uint64 addr;
            for (i = 0; i < NOFILE; i++) {
                if (p->vmaps[i] == 0) {
                    continue;
                }
                addr = (uint64) (p->vmaps[i]->addr);
                if (addr <= stval && stval < addr + p->vmaps[i]->len) {
                    vmap = p->vmaps[i];
                    break;
                }
            }
            if (i != NOFILE){
                char *mem = kalloc();
                int prot = PTE_U;
                if (mem == 0) {
                    p->killed = 1;
                } else {
                    memset(mem, 0, PGSIZE);
                    ilock(vmap->file->ip);
                }
            }
        }
    }
}

```

```

        readi(vmap->file->ip, 0, (uint64) mem, PGROUNDDOWN(stval -
addr), PGSIZE);
        iunlock(vmap->file->ip);
        if (vmap->prot & PROT_READ){
            prot |= PTE_R;
        }
        if (vmap->prot & PROT_WRITE){
            prot |= PTE_W;
        }
        if (mappages(p->pagetable, PGROUNDDOWN(stval), PGSIZE,
(uint64)mem, prot) != 0) {
            kfree(mem);
            p->killed = 1;
        }
    }
    }else {
        p->killed = 1;
    }
    } else {
        p->killed = 1;
    }
    }
} else {
    printf("usertrap(): unexpected scause %p pid=%d\n", r_scause(), p->pid);
    printf("                sepc=%p stval=%p\n", r_sepc(), r_stval());
    p->killed = 1;
}

if(p->killed)
    exit(-1);

// give up the CPU if this is a timer interrupt.
if(which_dev == 2)
    yield();

usertrapret();
}

```

exit.c

```

// Create a new process, copying the parent.
// Sets up child kernel stack to return as if from fork() system call.
int
fork(void)
{
    int i, pid;
    struct proc *np;
    struct proc *p = myproc();

    // Allocate process.
    if((np = allocproc()) == 0){
        return -1;
    }

    // Copy user memory from parent to child.
    if(uvmcopy(p->pagetable, np->pagetable, p->sz) < 0){
        freeproc(np);
    }
}

```

```

    release(&np->lock);
    return -1;
}

for(i=0; i < NOFILE; i++){
    if (p->vmmaps[i]){
        np->vmmaps[i] = vma_alloc();
        np->vmmaps[i]->addr = p->vmmaps[i]->addr;
        np->vmmaps[i]->len = p->vmmaps[i]->len;
        np->vmmaps[i]->prot = p->vmmaps[i]->prot;
        np->vmmaps[i]->flag = p->vmmaps[i]->flag;
        np->vmmaps[i]->file = p->vmmaps[i]->file;
        filedup(p->vmmaps[i]->file);
    }
}

np->sz = p->sz;

np->parent = p;

// copy saved user registers.
*(np->trapframe) = *(p->trapframe);

// Cause fork to return 0 in the child.
np->trapframe->a0 = 0;

// increment reference counts on open file descriptors.
for(i = 0; i < NOFILE; i++)
    if(p->ofile[i])
        np->ofile[i] = filedup(p->ofile[i]);
np->cwd = idup(p->cwd);

safestrcpy(np->name, p->name, sizeof(p->name));

pid = np->pid;

np->state = RUNNABLE;

release(&np->lock);

return pid;
}

// Exit the current process. Does not return.
// An exited process remains in the zombie state
// until its parent calls wait().
void
exit(int status)
{
    struct proc *p = myproc();

    if(p == initproc)
        panic("init exiting");

    for(int i = 0; i < NOFILE; i++){
        if(p->vmmaps[i]){
            struct vma *vmap = p->vmmaps[i];

```

```

        if (vmap->prot & PROT_WRITE && vmap->flag == MAP_SHARED){
            begin_op();
            ilock(vmap->file->ip);
            writei(vmap->file->ip, 1, (uint64)vmap->addr, 0, vmap->len);
            iunlock(vmap->file->ip);
            end_op();
        }
        fileclose(vmap->file);
        vma_free(vmap);
        p->vmmaps[i] = 0;
    }
}

// Close all open files.
for(int fd = 0; fd < NOFILE; fd++){
    if(p->ofile[fd]){
        struct file *f = p->ofile[fd];
        fileclose(f);
        p->ofile[fd] = 0;
    }
}

begin_op();
iput(p->cwd);
end_op();
p->cwd = 0;

// we might re-parent a child to init. we can't be precise about
// waking up init, since we can't acquire its lock once we've
// acquired any other proc lock. so wake up init whether that's
// necessary or not. init may miss this wakeup, but that seems
// harmless.
acquire(&initproc->lock);
wakeup1(initproc);
release(&initproc->lock);

// grab a copy of p->parent, to ensure that we unlock the same
// parent we locked. in case our parent gives us away to init while
// we're waiting for the parent lock. we may then race with an
// exiting parent, but the result will be a harmless spurious wakeup
// to a dead or wrong process; proc structs are never re-allocated
// as anything else.
acquire(&p->lock);
struct proc *original_parent = p->parent;
release(&p->lock);

// we need the parent's lock in order to wake it up from wait().
// the parent-then-child rule says we have to lock it first.
acquire(&original_parent->lock);

acquire(&p->lock);

// Give any children to init.
reparent(p);

// Parent might be sleeping in wait().
wakeup1(original_parent);

```



```
p->xstate = status;
p->state = ZOMBIE;

release(&original_parent->lock);

// Jump into the scheduler, never to return.
sched();
panic("zombie exit");
}
```

测试:

```
xv6 kernel is booting

hart 1 starting
hart 2 starting
init: starting sh
$ mmaptest
mmap_test starting
test mmap f
test mmap f: OK
test mmap private
test mmap private: OK
test mmap read-only
test mmap read-only: OK
test mmap read/write
test mmap read/write: OK
test mmap dirty
test mmap dirty: OK
test not-mapped unmap
test not-mapped unmap: OK
test mmap two files
test mmap two files: OK
mmap_test: ALL OK
fork_test starting
fork_test OK
mmaptest: all tests succeeded
```

```
$ usertests
usertests starting
test manywrites: OK
test execout: OK
test copyin: OK
test copyout: OK
test copyinstr1: OK
test copyinstr2: OK
test copyinstr3: OK
test rwsbrk: OK
test truncate1: OK
test truncate2: OK
test truncate3: OK
```

```
test exitwait: OK
test rmdot: OK
test fourteen: OK
test bigfile: OK
test dirfile: OK
test iref: OK
test forktest: OK
test bigdir: OK
ALL TESTS PASSED
```

可以通过 mmaptest 和 usertests 测试

3. 实验评分:

按照实验要求对本次实验的所有小实验进行评分, 结果如图:

```
== Test running mmaptest ==
$ make qemu-gdb
(5.0s)
== Test    mmaptest: mmap f ==
mmaptest: mmap f: OK
== Test    mmaptest: mmap private ==
mmaptest: mmap private: OK
== Test    mmaptest: mmap read-only ==
mmaptest: mmap read-only: OK
== Test    mmaptest: mmap read/write ==
mmaptest: mmap read/write: OK
== Test    mmaptest: mmap dirty ==
mmaptest: mmap dirty: OK
== Test    mmaptest: not-mapped unmap ==
mmaptest: not-mapped unmap: OK
== Test    mmaptest: two files ==
mmaptest: two files: OK
== Test    mmaptest: fork_test ==
mmaptest: fork_test: OK
== Test usertests ==
$ make qemu-gdb
usertests: OK (236.9s)
== Test time ==
time: OK
Score: 140/140
```

通过了所有测试

4. 问题以及解决办法:

a. 如何确定文件映射到用户虚拟地址空间的位置

本次实验是想实现文件到用户空间的映射, 其实也就是为文件指针设置一个合适的数据结构, 并将其放置在用户空间的一块区域上。但是文件应当被映射到具体哪个位置呢? 这里我采用的是将其映射到当前内存区域的最后面, 这样可以确保不会和已经存在的其他信息产生位置冲突或者覆盖。

5. 实验心得:

1. 了解了如何将文件映射到用户虚拟内存中，并将修改后内容写回磁盘
2. 进一步熟悉了懒加载和文件系统